

car_ad_EDA

December 31, 2020

1 Research on car sales ads

1.0.1 Author: Xia Cui

1.0.2 Introduction

The report presents the exploratory analysis of a dataset of 51525 free advertisements for vehicles on Crankshaft list, posted from 2018-05-01 to 2019-04-19.

At the data preprocessing stage, we will address the issues of missing values and wrong data types, and also add columns where it fits in order to maximize the information it can present.

The analysis that follows first of all examines the pattern of the advertisements in a range of factors such as price, model year, cylinders, and mileage. Then we will look at what are the characteristics of cars that tend to be gone quickly, as well as of those that stay on the list longer than most. Finally, we will take a closer look at car prices, and investigate which factors influence the price of a vehicle, using the two most popular vehicle types as examples.

1.0.3 Table of Content

Part 1: data importing

Part 2: data preprocessing

Part 3: adding new columns

Part 4: exploratory data analysis

Part 5: conclusion

```
[1]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
from IPython.display import display
import seaborn as sns
%matplotlib inline
```

1.0.4 Part 1: Data importing

```
[2]: vehicles = pd.read_csv('https://code.s3.yandex.net/datasets/vehicles_us.csv')
vehicles.head()
```

```
[2]:
```

	price	model_year	model	condition	cylinders	fuel	odometer	\
0	9400	2011.0	bmw x5	good	6.0	gas	145000.0	
1	25500	NaN	ford f-150	good	6.0	gas	88705.0	
2	5500	2013.0	hyundai sonata	like new	4.0	gas	110000.0	
3	1500	2003.0	ford f-150	fair	8.0	gas	NaN	
4	14900	2017.0	chrysler 200	excellent	4.0	gas	80903.0	

	transmission	type	paint_color	is_4wd	date_posted	days_listed
0	automatic	SUV	NaN	1.0	2018-06-23	19
1	automatic	pickup	white	1.0	2018-10-19	50
2	automatic	sedan	red	NaN	2019-02-07	79
3	automatic	pickup	NaN	NaN	2019-03-22	9
4	automatic	sedan	black	NaN	2019-04-02	28

```
[3]: vehicles.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 51525 entries, 0 to 51524
Data columns (total 13 columns):
#   Column          Non-Null Count  Dtype
---  -
0   price           51525 non-null  int64
1   model_year      47906 non-null  float64
2   model           51525 non-null  object
3   condition       51525 non-null  object
4   cylinders       46265 non-null  float64
5   fuel            51525 non-null  object
6   odometer        43633 non-null  float64
7   transmission    51525 non-null  object
8   type            51525 non-null  object
9   paint_color     42258 non-null  object
10  is_4wd          25572 non-null  float64
11  date_posted     51525 non-null  object
12  days_listed     51525 non-null  int64
dtypes: float64(4), int64(2), object(7)
memory usage: 5.1+ MB
```

```
[4]: vehicles.shape
```

```
[4]: (51525, 13)
```

Let's also take a look at the summary data of the numeric variables.

```
[5]: vehicles.describe()
```

```
[5]:
```

	price	model_year	cylinders	odometer	is_4wd \
count	51525.000000	47906.000000	46265.000000	43633.000000	25572.0
mean	12132.464920	2009.750470	6.125235	115553.461738	1.0
std	10040.803015	6.282065	1.660360	65094.611341	0.0
min	1.000000	1908.000000	3.000000	0.000000	1.0
25%	5000.000000	2006.000000	4.000000	70000.000000	1.0
50%	9000.000000	2011.000000	6.000000	113000.000000	1.0
75%	16839.000000	2014.000000	8.000000	155000.000000	1.0
max	375000.000000	2019.000000	12.000000	990000.000000	1.0

	days_listed
count	51525.000000
mean	39.55476
std	28.20427
min	0.000000
25%	19.000000
50%	33.000000
75%	53.000000
max	271.000000

1.0.5 Part 1 conclusion

The dataset contains 51525 rows and 13 columns. a quick look at its first 5 rows and the general information reveals a number of issues that need to be addressed in the data preprocessing stage next. These include missing values, data type (e.g. date_posted), and some suspicious values (e.g. the 1 dollar car price). There are also some extremely large values in, for example, car price, and odometers. These could potentially be outliers that might skew the data. We will further examine these in the next section.

1.0.6 Part 2: Data preprocessing

Addressing missing values

```
[6]: vehicles.isnull().sum(axis = 0)
```

```
[6]:
```

price	0
model_year	3619
model	0
condition	0
cylinders	5260
fuel	0
odometer	7892
transmission	0
type	0
paint_color	9267
is_4wd	25953

```

date_posted      0
days_listed     0
dtype: int64

```

We can see that, out of the 13 columns, 5 columns have missing values. To have a better idea of the impact of these, let's calculate the percentage of missing values in the dataset.

```
[7]: 100*vehicles.isnull().sum(axis = 0) / vehicles.shape[0]
```

```

[7]: price      0.000000
     model_year  7.023775
     model       0.000000
     condition   0.000000
     cylinders   10.208637
     fuel        0.000000
     odometer    15.316836
     transmission 0.000000
     type        0.000000
     paint_color 17.985444
     is_4wd      50.369723
     date_posted  0.000000
     days_listed  0.000000
     dtype: float64

```

Missing values for 'is_4wd'

It seems that none of the missing values can be considered trivial in the five columns, especially for the 'is_4wd' variable. However, a closer look at the data description earlier shows that all the values in this column is '1', indicating the car is 4_wd. It is reasonable to believe that the missing values should have been '0', indicating that the car is not 4_wd. So the missing values for this variable is relatively easier to solve. We can simply replace the NaN with '0'.

```

[8]: vehicles['is_4wd'] = vehicles['is_4wd'].fillna(0)
     vehicles['is_4wd'].value_counts()

```

```

[8]: 0.0    25953
     1.0    25572
     Name: is_4wd, dtype: int64

```

Before we decide what to do with the missing values in the other four columns, let's examine them first.

Missing values for 'model_year'

```
[9]: vehicles[vehicles['model_year'].isna()].head()
```

```

[9]:   price  model_year      model  condition  cylinders  fuel  odometer  \
1   25500         NaN  ford f-150      good         6.0  gas   88705.0
20   6990         NaN  chevrolet tahoe  excellent         8.0  gas  147485.0

```

65	12800	NaN	ford f-150	excellent	6.0	gas	108500.0
69	7800	NaN	ford f-150	like new	8.0	gas	97510.0
72	3650	NaN	subaru impreza	excellent	NaN	gas	74000.0

	transmission	type	paint_color	is_4wd	date_posted	days_listed
1	automatic	pickup	white	1.0	2018-10-19	50
20	automatic	SUV	silver	1.0	2018-08-05	28
65	automatic	pickup	white	0.0	2018-09-23	15
69	automatic	truck	white	1.0	2019-02-20	39
72	automatic	sedan	blue	1.0	2018-08-07	60

Is model_year dependant on model? let's have a look.

```
[19]: vehicles.groupby('model')['model_year'].value_counts()
```

```
[19]: model          model_year
      acura tl          2005.0          31
              2007.0          30
              2008.0          28
              2012.0          27
              2006.0          24
              ..
      volkswagen passat 2009.0          3
              1995.0          1
              1999.0          1
              2002.0          1
              2018.0          1
      Name: model_year, Length: 2226, dtype: int64
```

It might be tricky to replace the missing values in 'model_year'. Each model category has several different model years and is not really helpful in filling the missing values.

While it is reasonable to believe that, quite often, the more odometers a car has, the older it could be, to fill in the model year based on its odometer would be over generalizing. Moreover, model_year potentially could be a key parameter in its sale price, therefore we need to be especially careful how to replace these, or whether to replace these at all. For now, let's keep the column as it is and decide what to do with it later.

Missing value for 'cylinders'

```
[20]: vehicles[vehicles['cylinders'].isna()].head()
```

```
[20]:   price  model_year          model  condition  cylinders  fuel  odometer \
9    9200    2008.0    honda pilot  excellent        NaN  gas  147191.0
36  10499    2013.0  chrysler 300      good        NaN  gas   88042.0
37   7500    2005.0  toyota tacoma      good        NaN  gas  160000.0
59   5200    2006.0  toyota highlander      good        NaN  gas  186000.0
63  30000    1966.0   ford mustang  excellent        NaN  gas   51000.0
```

	transmission	type	paint_color	is_4wd	date_posted	days_listed
9	automatic	SUV	blue	1.0	2019-02-15	17
36	automatic	sedan	NaN	0.0	2018-05-05	22
37	automatic	pickup	NaN	0.0	2018-07-22	44
59	automatic	SUV	green	0.0	2018-12-20	2
63	manual	convertible	red	0.0	2019-01-23	17

It is possible that the number of cylinders is dependant on the car model. Let's check using the first two in dataset above.

```
[25]: vehicles[vehicles['model']=='honda pilot']['cylinders'].value_counts()
```

```
[25]: 6.0    269
      4.0     2
      8.0     2
      5.0     1
      Name: cylinders, dtype: int64
```

```
[26]: vehicles[vehicles['model']=='chrysler 300']['cylinders'].value_counts()
```

```
[26]: 6.0    216
      8.0    58
      4.0     5
      Name: cylinders, dtype: int64
```

Our assumption seems to hold some truth. For each car model, there is a dominant cylinder number. Although it would not be entirely accurate, it's relatively safe to use the dominant cylinder number of each model to replace the missing values in the cylinder column.

```
[27]: vehicles["cylinders"] = vehicles.groupby("model")["cylinders"].transform(lambda x: x.fillna(x.mode()[0]))
```

Let's check if the null values have been filled.

```
[28]: vehicles["cylinders"].isna().sum()
```

```
[28]: 0
```

It's all done for cylinders. Now let's continue with 'odometer'.

Missing values for 'odometer'

```
[335]: vehicles[vehicles['odometer'].isna()].head()
```

```
[335]:   price  model_year      model  condition  cylinders  fuel  odometer \
3   1500    2003.0  ford f-150    fair        8.0  gas      NaN
15  17990    2013.0    ram 1500  excellent    8.0  gas      NaN
23   7500    2004.0  jeep wrangler    good        6.0  gas      NaN
24   3950    2009.0  chrysler 200  excellent    4.0  gas      NaN
```

25	11499	2017.0	chevrolet malibu	like new	4.0	gas	NaN
----	-------	--------	------------------	----------	-----	-----	-----

	transmission	type	paint_color	is_4wd	date_posted	days_listed
3	automatic	pickup	NaN	0.0	2019-03-22	9
15	automatic	pickup	red	1.0	2018-05-15	111
23	automatic	SUV	red	1.0	2018-05-17	39
24	automatic	sedan	red	0.0	2018-06-11	40
25	automatic	sedan	NaN	0.0	2018-07-26	43

It is possible to fill up the missing values in odometer using the median values based on 'model_year' and 'condition'. However, we know that there are missing values in 'model_year' as well. So we need to check how many have missing values in both columns.

```
[29]: vehicles[(vehicles['odometer'].isna()) & (vehicles['model_year'].isna())]
```

```
[29]:
```

	price	model_year	model	condition	cylinders	\
159	23300	NaN	nissan frontier crew cab sv	good	6.0	
260	14975	NaN	toyota 4runner	good	6.0	
370	4700	NaN	kia soul	good	4.0	
586	26000	NaN	toyota rav4	like new	4.0	
659	8400	NaN	volkswagen jetta	good	4.0	
...	
51195	21999	NaN	ram 2500	good	6.0	
51222	1000	NaN	acura tl	good	6.0	
51257	6500	NaN	toyota corolla	good	4.0	
51295	3850	NaN	hyundai elantra	excellent	4.0	
51399	4400	NaN	kia sorento	excellent	6.0	

	fuel	odometer	transmission	type	paint_color	is_4wd	date_posted	\
159	gas	NaN	other	pickup	grey	1.0	2018-07-24	
260	gas	NaN	automatic	SUV	silver	0.0	2018-05-13	
370	gas	NaN	manual	sedan	white	0.0	2019-01-14	
586	gas	NaN	automatic	SUV	NaN	0.0	2018-08-09	
659	diesel	NaN	manual	wagon	NaN	0.0	2018-10-22	

	fuel	odometer	transmission	type	paint_color	is_4wd	date_posted	\
...	
51195	diesel	NaN	automatic	truck	white	1.0	2018-05-10	
51222	gas	NaN	automatic	sedan	grey	0.0	2018-12-09	
51257	gas	NaN	automatic	sedan	white	0.0	2018-10-16	
51295	gas	NaN	automatic	sedan	silver	0.0	2019-03-16	
51399	gas	NaN	automatic	SUV	silver	0.0	2018-08-21	

	days_listed
159	73
260	57
370	50
586	29
659	37

```
...
51195      35
51222      23
51257      75
51295      83
51399      23
```

[549 rows x 13 columns]

There are 549 rows that have missing values in both ‘model_year’ and ‘odometer’. For these 549 rows, we can use the median odometers by condition only to fill in the missing values. For the rest, we can use both ‘model_year’ and ‘condition’ to do so.

```
[31]: vehicles.loc[vehicles['model_year'].notna(), "odometer"] = vehicles[
        vehicles['model_year'].notna()].groupby(["model_year", "condition"])[
        "odometer"].transform(lambda x: x.fillna(x.median()))
```

```
[32]: vehicles.loc[vehicles['model_year'].isna(), "odometer"] = vehicles[
        vehicles['model_year'].isna()].groupby("condition")[
        "odometer"].transform(lambda x: x.fillna(x.median()))
```

Now let’s check how many missing values are left in the ‘odometer’ column.

```
[34]: vehicles['odometer'].isna().sum()
```

[34]: 7

Let’s found out who these are.

```
[35]: vehicles[vehicles['odometer'].isna()]
```

```
[35]:
```

	price	model_year	model	condition	cylinders	fuel	\
21421	4500	1974.0	chevrolet corvette	fair	8.0	gas	
28009	65000	1960.0	chevrolet corvette	like new	8.0	gas	
31806	1700	1996.0	ford mustang	salvage	6.0	gas	
33257	4500	1963.0	chevrolet impala	fair	6.0	gas	
33907	12995	1908.0	cadillac escalade	excellent	8.0	gas	
45694	18000	1929.0	ford f-150	good	8.0	gas	
46911	22300	2003.0	chevrolet corvette	new	8.0	gas	

	odometer	transmission	type	paint_color	is_4wd	date_posted	\
21421	NaN	automatic	sedan	red	0.0	2018-12-15	
28009	NaN	manual	coupe	NaN	0.0	2018-11-03	
31806	NaN	manual	convertible	white	0.0	2019-03-31	
33257	NaN	automatic	sedan	NaN	0.0	2019-03-17	
33907	NaN	automatic	SUV	white	0.0	2018-06-24	
45694	NaN	manual	other	silver	0.0	2018-11-18	
46911	NaN	manual	convertible	black	0.0	2018-11-08	

	days_listed
21421	18
28009	41
31806	46
33257	38
33907	25
45694	59
46911	23

These cars have both model year and conditions, but why aren't their values filled? Let's take a look at the first year condition combo.

```
[36]: vehicles[(vehicles['model_year'] == 1974.0) & (vehicles['model'] == 'chevrolet_
      ↪corvette')]
```

```
[36]:
```

	price	model_year	model	condition	cylinders	fuel	\
4708	10500	1974.0	chevrolet corvette	good	8.0	gas	
16723	4950	1974.0	chevrolet corvette	good	8.0	gas	
21421	4500	1974.0	chevrolet corvette	fair	8.0	gas	

	odometer	transmission	type	paint_color	is_4wd	date_posted	\
4708	4133.0	automatic	coupe	red	0.0	2018-09-16	
16723	29000.0	automatic	coupe	blue	0.0	2018-07-21	
21421	NaN	automatic	sedan	red	0.0	2018-12-15	

	days_listed
4708	74
16723	103
21421	18

The result above shows that the reason the missing odometer is not filled is because the other two are both in good conditions and there is no median value for the fair condition car from the same year.

There are only 7 left. We can leave these for now.

Missing values in 'paint_color'

```
[37]: vehicles[vehicles['paint_color'].isna()].head()
```

```
[37]:
```

	price	model_year	model	condition	cylinders	fuel	odometer	\
0	9400	2011.0	bmw x5	good	6.0	gas	145000.0	
3	1500	2003.0	ford f-150	fair	8.0	gas	193850.0	
8	11500	2012.0	kia sorento	excellent	4.0	gas	104174.0	
12	18990	2012.0	ram 1500	excellent	8.0	gas	140742.0	
21	5250	2007.0	toyota rav4	good	6.0	gas	154000.0	

	transmission	type	paint_color	is_4wd	date_posted	days_listed
--	--------------	------	-------------	--------	-------------	-------------

0	automatic	SUV	NaN	1.0	2018-06-23	19
3	automatic	pickup	NaN	0.0	2019-03-22	9
8	automatic	SUV	NaN	1.0	2018-07-16	19
12	automatic	pickup	NaN	1.0	2019-04-02	37
21	automatic	SUV	NaN	0.0	2018-08-22	8

Let's first have a look at how the colors are distributed.

```
[38]: vehicles['paint_color'].value_counts()
```

```
[38]: white      10029
      black      7692
      silver     6244
      grey       5037
      blue       4475
      red        4421
      green      1396
      brown      1223
      custom     1153
      yellow      255
      orange      231
      purple      102
      Name: paint_color, dtype: int64
```

We can either fill the missing values with the dominant car color, 'white', or with the dominate car color of same model. Let's take the second approach so white wouldn't become even more dominant.

```
[39]: vehicles["paint_color"] = vehicles.groupby("model")["paint_color"].transform(lambda x: x.fillna(x.mode()[0]))
```

```
[40]: vehicles['paint_color'].isna().sum()
```

```
[40]: 0
```

Now we've filled up all the NA values that we could. Let's check what's there left.

```
[41]: vehicles.isnull().sum(axis = 0)
```

```
[41]: price          0
      model_year    3619
      model         0
      condition     0
      cylinders     0
      fuel          0
      odometer      7
      transmission  0
      type          0
```

```
paint_color      0
is_4wd           0
date_posted      0
days_listed     0
dtype: int64
```

While we are here, let's also check if there are any duplicated rows.

```
[42]: vehicles.duplicated().sum()
```

```
[42]: 0
```

No duplicated row!!

All the missing values left are the 'model_year' missing values, and those 'odometer' missing values where the 'model_year' is also missing. 3619 rows represent over 7% of the total dataset and is not insignificant. However, 'model_year' is a key parameter for car price and keeping these in the dataset or replace them with any other values will potentially skew the data and impact the results. let's drop these and save the dataset with a new name.

Drop NA values

```
[43]: vehicles_new = vehicles.dropna().reset_index(drop = True)
vehicles_new.shape
```

```
[43]: (47899, 13)
```

The new dataset has 47899 rows and 13 columns. Let's double check to see if we missed anything.

```
[45]: vehicles_new.isnull().sum()
```

```
[45]: price      0
model_year    0
model         0
condition     0
cylinders     0
fuel          0
odometer      0
transmission  0
type          0
paint_color   0
is_4wd        0
date_posted   0
days_listed  0
dtype: int64
```

Now there are no NA values in the dataset. We can move on to change the data types and add new columns. First, let's have another look at the datatypes. **Changing data types**

```
[46]: vehicles_new.dtypes
```

```
[46]: price            int64
      model_year      float64
      model           object
      condition       object
      cylinders       float64
      fuel            object
      odometer        float64
      transmission    object
      type            object
      paint_color     object
      is_4wd          float64
      date_posted     object
      days_listed     int64
      dtype: object
```

'date_posted' should be datetime data type in order to calculate the day of the week, month and year later on.

cylinders, and is_4wd should be integer type.

'model_year' is float. We can leave it be because later the extracted year of the ad posting date will be integers and we can still calculate the age of the cars using those two columns.

```
[47]: vehicles_new['model_year'] = vehicles_new['model_year'].astype(int)
      vehicles_new.head()
```

```
[47]:   price  model_year      model  condition  cylinders  fuel  odometer  \
0   9400         2011    bmw x5      good        6.0  gas   145000.0
1   5500         2013  hyundai sonata  like new        4.0  gas   110000.0
2   1500         2003    ford f-150    fair        8.0  gas   193850.0
3  14900         2017  chrysler 200  excellent        4.0  gas    80903.0
4  14990         2014  chrysler 300  excellent        6.0  gas    57954.0

      transmission  type  paint_color  is_4wd  date_posted  days_listed
0   automatic    SUV      black      1.0   2018-06-23          19
1   automatic    sedan      red      0.0   2019-02-07          79
2   automatic  pickup    white      0.0   2019-03-22           9
3   automatic    sedan    black      0.0   2019-04-02          28
4   automatic    sedan    black      1.0   2018-06-20          15
```

```
[48]: vehicles_new['date_posted'] = pd.to_datetime(vehicles_new['date_posted'],
      ↪format = '%Y-%m-%d')
```

```
[49]: vehicles_new['cylinders'] = vehicles_new['cylinders'].astype(int)
```

```
[50]: vehicles_new['is_4wd'] = vehicles_new['is_4wd'].astype(int)
```

```
[51]: vehicles_new.dtypes
```

```
[51]: price                int64
      model_year          int64
      model               object
      condition           object
      cylinders           int64
      fuel               object
      odometer            float64
      transmission       object
      type               object
      paint_color         object
      is_4wd             int64
      date_posted         datetime64[ns]
      days_listed         int64
      dtype: object
```

Now we have the data in the types exactly what we want. Let's move on to add some new columns.

1.0.7 Part 3 : Make calculations and add them to the table

Add a column of the day of the week when the ads were posted

```
[52]: vehicles_new['day_of_week_posted'] = vehicles_new['date_posted'].dt.weekday
```

Add column of the month the ads were placed

```
[53]: vehicles_new['month_posted'] = pd.DatetimeIndex(vehicles_new['date_posted']).
      ↪month
```

Add a column of the year the ads were posted

```
[54]: vehicles_new['year_posted'] = pd.DatetimeIndex(vehicles_new['date_posted']).year
```

Add a column of the vehicle's age (in years) when the ads were placed

```
[55]: vehicles_new['vehicle_age'] = vehicles_new['year_posted'] -
      ↪vehicles_new['model_year']
```

Add a column of the vehicle's average mileage per year

```
[56]: vehicles_new['mileage_per_year'] = vehicles_new['odometer'] /
      ↪vehicles_new['vehicle_age']
```

In the condition column, replace string values with a numeric scale

```
[57]: vehicles_new['condition'] = vehicles_new[
      ↪'condition'].replace(['salvage', 'fair', 'good', 'excellent', 'like new'],
      ↪'new'),
```

```
[0, 1, 2, 3, 4, 5])
vehicles_new.head()
```

```
[57]:   price  model_year      model  condition  cylinders fuel  odometer  \
0   9400        2011      bmw x5          2          6  gas  145000.0
1   5500        2013  hyundai sonata        4          4  gas  110000.0
2   1500        2003    ford f-150         1          8  gas  193850.0
3  14900        2017  chrysler 200         3          4  gas   80903.0
4  14990        2014  chrysler 300         3          6  gas   57954.0

   transmission  type  paint_color  is_4wd  date_posted  days_listed  \
0   automatic    SUV      black      1  2018-06-23          19
1   automatic    sedan      red      0  2019-02-07          79
2   automatic  pickup    white      0  2019-03-22           9
3   automatic    sedan    black      0  2019-04-02          28
4   automatic    sedan    black      1  2018-06-20          15

   day_of_week_posted  month_posted  year_posted  vehicle_age  \
0                   5              6          2018           7
1                   3              2          2019           6
2                   4              3          2019          16
3                   1              4          2019           2
4                   2              6          2018           4

   mileage_per_year
0      20714.285714
1      18333.333333
2      12115.625000
3      40451.500000
4      14488.500000
```

The `mileage_per_year` column has many digits after the decimal point. Let's round these to the nearest whole number up.

```
[59]: vehicles_new['mileage_per_year'] = vehicles_new['mileage_per_year'].apply(np.
      ↪ ceil)
```

1.0.8 Part 3 Conclusion

Now the dataset is more telling, including which weekday, month, and year the advertisements were placed, as well as vehicle's age and their mileage per year. We can easily draw on these addition columns in the exploratory analysis next.

1.0.9 Part 4: Exploratory data analysis

The exploratory analysis will first examine the following parameters: price, vehicle's age when the ad was placed, mileage, number of cylinders, and condition. A histogram or bar graph will be plotted for each of these parameters.

Such informaton, together with the descriptive statistics, will allow us to identify outliers in the data, and understand how they might affect the form and readability of the graphs. Then we will decide what outliers to be filtered out of the data, and plot another set of graphs to show the data distribution along the above mentioned parameters.

Then, we will have a look at how many days advertisements were displayed. The histogram and descriptive statistics of this parameter will allow us to understand the typical lifetime of an ad.

Last, we'll analyze the number of ads and the average price for each type of vehicle. The two types with the greatest number of ads will then be further examined to find out which factors impact on their car prices most.

4.1: Distribution of car price, age, mileage, cylinders, and condition

4.2: Removing outliers

4.3: Plotting new graphs

4.4: Studying the lifespan of advertisements


4.5: Number of ads and mean prices by car type

4.6: Factors in car prices**

1.0.10 4.1 Distribution of car price, age, mileage, cylinders, and condition

Price distribution

```
[108]: vehicles_new.hist('price', bins = 100, grid = True, figsize = (8, 4))
plt.xlabel('Price')
plt.ylabel('Frequency')
plt.title('Price histogram')
plt.show()
```



```
[61]: vehicles_new['price'].describe()
```

```
[61]: count      47899.000000
      mean       12159.549281
      std        10079.922175
      min         1.000000
      25%        5000.000000
      50%        9000.000000
      75%       16900.000000
      max       375000.000000
      Name: price, dtype: float64
```

Both the histogram and statistics data of the 'price' column shows that there are outliers on both sides of the distribution. It's clear that 75% of the cars are priced 16900 and under, and the number of cars priced over 75000 is approaching zero. However, it's not clear what the distribution is like on the cheaper end. To find out, let's have a look at the 10 cheapest cars.

```
[62]: vehicles_new.sort_values('price').head(10)
```

```
[62]:      price  model_year      model  condition  cylinders  fuel  \
10951      1         2015      ram 1500          3          10  gas
8804       1         2015    ford edge          3           6  gas
11438      1         2014  chevrolet silverado        4           8  gas
38692      1         2002    volkswagen jetta        2           4  gas
8702       1         2014      gmc sierra          3           8  gas
8664       1         2014    chevrolet camaro        3           6  gas
8663       1         2016      jeep wrangler        3           6  gas
46210      1         2007  chevrolet trailblazer        3           8  gas
8662       1         2012      ford mustang          3          10  gas
8661       1         2017      ram 3500          3          10  gas
```

```
      odometer  transmission  type  paint_color  is_4wd  date_posted  \
10951   59980.0         other  truck      silver        1  2018-07-25
8804   24897.0      automatic   SUV      black        1  2018-06-21
11438    42.0      automatic  truck      black        1  2018-06-28
38692  167062.0      automatic  sedan      blue        0  2019-02-08
8702  147470.0      automatic  truck      brown        0  2018-12-29
8664   51550.0        manual  coupe      black        1  2018-07-31
8663   56000.0      automatic   SUV        red        1  2018-05-20
46210  137000.0      automatic   SUV      black        1  2018-08-06
8662   41469.0      automatic  coupe      orange        1  2018-08-13
8661   57482.0         other  truck      white        1  2019-03-17
```

```
      days_listed  day_of_week_posted  month_posted  year_posted  \
10951           4                   2             7          2018
```


8804	83	3	6	2018
11438	60	3	6	2018
38692	12	4	2	2019
8702	9	5	12	2018
8664	36	1	7	2018
8663	44	6	5	2018
46210	28	0	8	2018
8662	20	0	8	2018
8661	3	6	3	2019

	vehicle_age	mileage_per_year
10951	3	19994.0
8804	3	8299.0
11438	4	11.0
38692	17	9828.0
8702	4	36868.0
8664	4	12888.0
8663	2	28000.0
46210	11	12455.0
8662	6	6912.0
8661	2	28741.0

It's interesting that there are so many 1 dollar cars, and these are not particularly old cars or cars have high mileage. Could that be a mistake when entering the data? Can we safely keep these out? Let's find out the proportion of these cars.

```
[63]: len(vehicles_new.query('price == 1')) / len(vehicles_new)
```

```
[63]: 0.015574437879705213
```

The cars priced exactly 1 dollar represents slightly over 1 percent of the total cars. Therefore we can consider it safe to remove these outliers on the cheaper end of the price range. Now. Let's have a look at how many cars more priced over 50000, and 75000 respectively.

```
[64]: len(vehicles_new.query('price >= 50000')) / len(vehicles_new)
```

```
[64]: 0.004551243240986242
```


```
[65]: len(vehicles_new.query('price >=75000')) / len(vehicles_new)
```

```
[65]: 0.0005219315643332846
```

They both represent a very small percentage of the total cars. While it would be interesting to study these outliers, they also skew the data towards the more expensive end. For now, let's arbitrarily use 50000 as the upper threshold for the data. Please note that this is still higher than the upper whisker (3rd quantile + 1.5 * IQR), which is 34750 if we used a boxplot to show the price distribution.

Vehicle age distribution

```
[110]: vehicles_new.hist('vehicle_age', bins = 100, grid = True, figsize = (8,4))
plt.xlabel('Vehicle age')
plt.ylabel('Frequency')
plt.title('Car age histogram')
plt.show()
```



```
[67]: vehicles_new['vehicle_age'].describe()
```

```
[67]: count      47899.000000
      mean         8.549970
      std         6.257531
      min          0.000000
      25%          4.000000
      50%          7.000000
      75%         12.000000
      max         110.000000
      Name: vehicle_age, dtype: float64
```

The distribution of the 'vehicle_age' shows two peaks, one at 0, meaning there are lots of cars on sale which are quite new, and one around 10 years. 75% of the cars are aged 12 years and under, but there are some outliers on the older side of the distribution, the oldest one being 110 years old. The distribution seems to be thinning out to 0 after 40 years. Let's have a look at the proportion of cars older than 40 years.

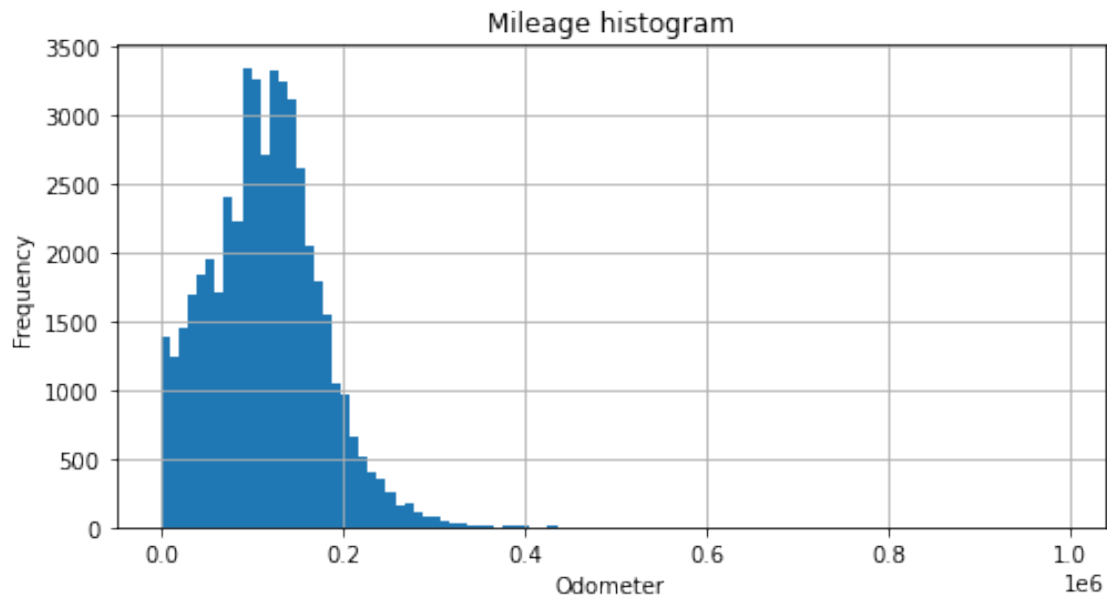
```
[68]: len(vehicles_new.query('vehicle_age > 40'))/len(vehicles_new)
```

```
[68]: 0.0032986074865863587
```

Cars older than 40 years only represents a very small proportion. It's safe to remove these.

Odometer distribution

```
[111]: vehicles_new.hist('odometer', bins = 100, grid = True, figsize = (8, 4))
plt.xlabel('Odometer')
plt.ylabel('Frequency')
plt.title('Mileage histogram')
plt.show()
```



```
[112]: vehicles_new['odometer'].describe()
```

```
[112]: count      47899.000000
mean       115081.371135
std        62424.773267
min         0.000000
25%        72757.000000
50%       114524.000000
75%       152869.000000
max       990000.000000
Name: odometer, dtype: float64
```

As show in the graph and descriptive statistics, 75% of the cars have odometers less than 153000. After 300000, the numbers is approaching zero. There are clearly some outliers on the high odometers' end. Let's find out the proportion of cars having over 300000 odometers (3rd quantile + 1.5 * IQR is 273000).

```
[113]: len(vehicles_new.query('odometer > 300000')) / len(vehicles_new)
```

```
[113]: 0.005490720056786154
```

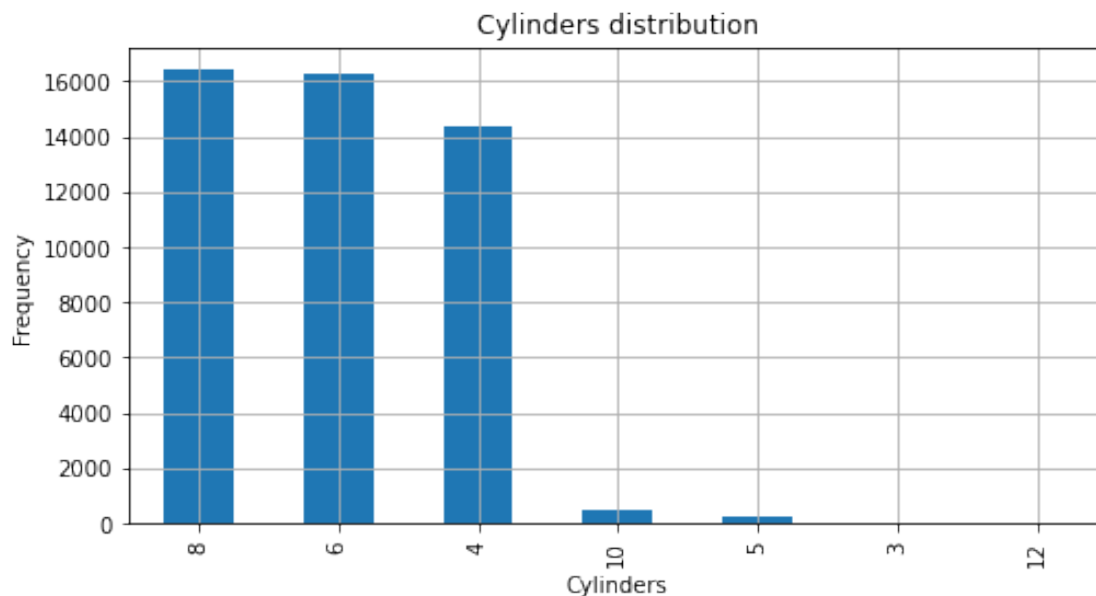
```
[114]: vehicles_new.query('odometer > 300000')['vehicle_age'].describe()
```

```
[114]: count    263.000000  
      mean     15.433460  
      std      6.102366  
      min      0.000000  
      25%     12.000000  
      50%     15.000000  
      75%     19.000000  
      max     54.000000  
      Name: vehicle_age, dtype: float64
```

It's a very small proportion. A quick look at the age distribution of these high odometer cars also shows that these are not exclusively old cars only. We can consider removing these outliers.

Cylinders

```
[115]: vehicles_new['cylinders'].value_counts().plot(kind = 'bar', grid = True,  
      ↳figsize = (8, 4))  
plt.xlabel('Cylinders')  
plt.ylabel('Frequency')  
plt.title('Cylinders distribution')  
plt.show()
```



As shown in the graph, most of the cars have 4, 6, or 8 cylinders. The bar for 3 and 12 is very low and it's really hard to see. Let's find out how many cars have 3 or 12 cylinders.

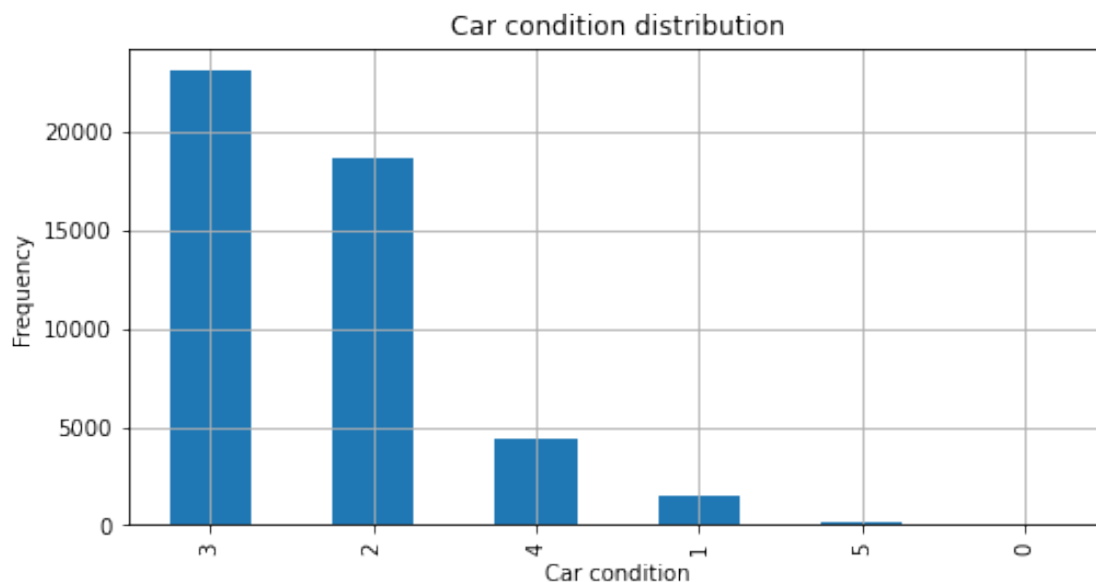
```
[116]: len(vehicles_new.query('cylinders in (3, 12)'))
```

```
[116]: 35
```

Only 35!

Car conditions

```
[117]: vehicles_new['condition'].value_counts().plot(kind = 'bar',grid = True, figsize_
      ↪= (8, 4))
plt.xlabel('Car condition')
plt.ylabel('Frequency')
plt.title('Car condition distribution')
plt.show()
```



As shown above, most of the cars' condition are 3, or 2, excellent and good respectively. There are some cars that are new (5), and barely any are salvage(0)

1.0.11 4.2 Removing outliers

Given the information shown in the graphs above, let's remove the cars that are priced 1 or over 50000, older than 40 years, and have odometers of more than 300000. In order not to lose any data, we will keep the outliers in a separate dataset.

Slicing data

```
[118]: vehicles_filtered = vehicles_new.query(
      'price != 1 & price <= 50000 & vehicle_age <= 40 & odometer <= 300000')
```

```
[119]: len(vehicles_filtered)/len(vehicles_new)
```

```
[119]: 0.971398150274536
```

Slightly less than 3 percent of the data is filtered out. It's reasonable. We can proceed to save the data that is filtered out in a separate dataset and use the filtered data for the following exploratory analysis.

```
[120]: vehicles_filtered_out = vehicles_new.query(  
        'price == 1 | price > 50000 | vehicle_age > 40 | odometer > 300000')
```

```
[121]: len(vehicles_filtered_out)/len(vehicles_new)
```


```
[121]: 0.028601849725463997
```

1.0.12 4.3 New graphs

Let's use the filtered data to plot new histograms. To see the comparison, we'll plot the two graphs next to each other.

Car price

```
[124]: fig, (ax1,ax2) = plt.subplots(1,2, figsize=(10,4))  
vehicles_new['price'].hist(bins = 50, grid = True, ax = ax1)  
ax1.set_title('Full data')  
ax1.set_ylabel('Frequency')  
ax1.set_xlabel('Car price')  
vehicles_filtered['price'].hist(bins = 50, grid = True, color = 'orange',ax =  
    ↪ax2)  
ax2.set_title('Filtered data')  
ax2.set_xlabel('Car price')  
  
plt.tight_layout()
```




Now the histogram for the filtered data is more readable. It shows the peak of the car prices just under 5000, and then the number decreases as the price goes up.

Let's plot the histograms for the vehicles' age next.

Vehicle age

```
[125]: fig, (ax1,ax2) = plt.subplots(1,2, figsize=(10,4))
vehicles_new['vehicle_age'].hist(bins = 50, grid = True, ax = ax1)
ax1.set_title('Full data')
ax1.set_ylabel('Frequency')
ax1.set_xlabel('Vehicle age')

vehicles_filtered['vehicle_age'].hist(bins = 50, grid = True, color = 'orange',
→ax = ax2)
ax2.set_title('Filtered data')
ax2.set_xlabel('Vehicle age')
plt.tight_layout()
```



The full data and filtered data exhibit different peaks. While there are 3 peaks for the full data, close 0, 8, and 12 year, the peak for the filtered data center around 4-7 years old. Did the removal of these 1 dollar cars or more expensive cars do that to the data? Let's have a look.

```
[126]: vehicles_filtered_out['vehicle_age'].describe()
```


```
[126]: count      1370.000000
mean         10.332117
std          15.875011
min           0.000000
25%           0.000000
50%           3.000000
75%          14.000000
max          110.000000
```

Name: vehicle_age, dtype: float64

Indeed, in the filtered out data, the mean age of the cars is 10 years old. Taking these away from the data, the vehicle age distribution shows different peaks from before. We are not going to do anything for the moment but it's good to keep this in mind.

Odometer

```
[129]: fig, (ax1,ax2) = plt.subplots(1,2, figsize=(10,4))
vehicles_new['odometer'].hist(bins = 50, grid = True, ax = ax1)
ax1.set_title('Full data')
ax1.set_ylabel('Frequency')
ax1.set_xlabel('Odometers')
vehicles_filtered['odometer'].hist(bins = 25, grid = True, color = 'orange', ax=
    ↪ ax2)
ax2.set_title('Filtered data')
ax2.set_xlabel('Odometers')
plt.tight_layout()
```



The histogram of the mileage from the filtered data is much more readable and informative. It shows a peak around 125000, and right skewed.


Cylinders

```
[130]: fig, (ax1,ax2) = plt.subplots(1,2, figsize=(10,4))
vehicles_new['cylinders'].value_counts().plot(kind = 'bar', grid = True, ax =
    ↪ ax1)
ax1.set_title('Full data')
ax1.set_ylabel('Frequency')
ax1.set_xlabel('Cylinders')

vehicles_filtered['cylinders'].value_counts().plot(
    kind = 'bar', grid = True, color = 'orange', ax = ax2)
```



```
ax2.set_title('Filtered data')
ax2.set_xlabel('Cylinders')
plt.tight_layout()
```




Filtering the data changed the ranking position of 6 and 8, as well as 5 and 10. This indicates that most of the cars filtered out have 8 cylinders.

Car condition

```
[131]: fig, (ax1,ax2) = plt.subplots(1,2, figsize=(10,4))
vehicles_new['condition'].value_counts().plot(kind = 'bar', grid = True, ax = ax1)
ax1.set_title('Full data')
ax1.set_ylabel('Frequency')
ax1.set_xlabel('Car conditions')

vehicles_filtered['condition'].value_counts().plot(
    kind = 'bar', grid = True, color = 'orange', ax = ax2)
ax2.set_title('Filtered data')
ax2.set_xlabel('Car conditions')
plt.tight_layout()
```




The distribution of the car conditions didn't change much, which is great! The cars that dominate are still **good** and **excellent**!

1.0.13 4.4 Life span of car advertisements

In this section, we will study how many days advertisements were displayed (`days_listed`) and describe the typical lifetime of an ad. We will also take a closer look at what kind of cars were removed quickly, and what kind were listed for an abnormally long time.

Just to be sure that, by filtering the data, we didn't change the distribution of this parameter dramatically, let's plot two histograms for '`days_listed`', using the full dataset, and the filtered dataset.

```
[132]: vehicles_new['days_listed'].hist(bins = 50, figsize = (8, 4))
vehicles_filtered['days_listed'].hist(bins = 50, figsize = (8, 4))
plt.xlabel('Number of days the advertisement were listed')
plt.ylabel('Frequency')
plt.show()
```



Filtering the data didn't seem to change the distribution of 'days_listed' much, which is good news.

```
[133]: vehicles_filtered['days_listed'].describe()
```

```
[133]: count    46529.000000
      mean      39.571901
      std       28.200632
      min        0.000000
      25%       19.000000
      50%       33.000000
      75%       53.000000
      max      271.000000
      Name: days_listed, dtype: float64
```

Using the filtered dataset, the mean for days listed is approximately 39 days, and the median 33 days. 75% of the ads were listed 53 days and under. This indicates that the typical life of a car ad is slightly over a month but less than 2 months.

Taking a look at the extreme ends of the ad life: the longest ad life is 271 days, whereas there seems to be a lot of ads which lasted less than one day! From the graph we can also see that the numbers start to thin out fast after 150.

Let's find out which cars have the days listed as 0, which ones have less than 7 days, which ones were listed over 150 days, and which car has the longest listing days.

Zero day listings

```
[134]: vehicles_filtered.query('days_listed == 0').head()
```

```
[134]:
```

	price	model_year	model	condition	cylinders	\
1165	14995	2008	chevrolet silverado 1500	3	8	
1854	14000	1999	ford f250	3	8	
2707	4000	2004	ram 1500	3	8	
3702	16750	1985	chevrolet corvette	4	8	
4254	5000	2007	toyota corolla	2	4	

	fuel	odometer	transmission	type	paint_color	is_4wd	\
1165	gas	93300.0	automatic	truck	grey	1	
1854	diesel	137500.0	automatic	truck	red	1	
2707	gas	250000.0	automatic	truck	brown	1	
3702	gas	24540.0	automatic	hatchback	white	0	
4254	gas	223000.0	manual	sedan	silver	0	

	date_posted	days_listed	day_of_week_posted	month_posted	year_posted	\
1165	2018-05-15	0	1	5	2018	
1854	2018-09-27	0	3	9	2018	
2707	2018-08-13	0	0	8	2018	
3702	2018-10-14	0	6	10	2018	
4254	2018-07-11	0	2	7	2018	

	vehicle_age	mileage_per_year
1165	10	9330.0
1854	19	7237.0
2707	14	17858.0
3702	33	744.0
4254	11	20273.0

```
[135]: len(vehicles_filtered.query('days_listed == 0'))
```

```
[135]: 49
```

There doesn't seem to be any particular pattern among the car ads that didn't last even for a day. Could this be a mistake?

In a real life situation, this would need to be brought to attention to the team who provided the data for further information. For the purpose of this report, let's have a look at ads that lasted more than 0 but gone within a week. We will arbitrarily decide these are the ones that are gone quickly.

Cars with short advertisement life

```
[136]: vehicles_filtered.query('0 < days_listed <= 7')
```

```
[136]:
```

	price	model_year	model	condition	cylinders	\
30	9499	2015	nissan altima	4	4	
37	8000	2009	ford f-150	2	8	
57	5200	2006	toyota highlander	2	6	
70	6950	2005	chevrolet tahoe	3	8	

111	33900	2018	chevrolet silverado 1500 crew	2	8
...
47792	12990	2013	honda accord	3	6
47830	4700	2007	toyota corolla	3	4
47858	7300	2016	ford fusion se	3	4
47870	9500	2012	chevrolet traverse	2	6
47879	20481	2018	toyota camry	4	4

	fuel	odometer	transmission	type	paint_color	is_4wd	date_posted	\
30	gas	51848.0	automatic	sedan	grey	0	2018-11-12	
37	gas	234000.0	automatic	truck	black	1	2019-03-31	
57	gas	186000.0	automatic	SUV	green	0	2018-12-20	
70	gas	186021.0	automatic	SUV	black	1	2018-10-30	
111	gas	11315.0	other	pickup	white	1	2019-03-01	
...	
47792	gas	118659.0	automatic	coupe	red	0	2018-05-02	
47830	gas	137000.0	automatic	sedan	grey	0	2018-07-17	
47858	gas	106212.0	automatic	sedan	grey	0	2019-03-10	
47870	gas	144500.0	automatic	SUV	silver	1	2019-03-05	
47879	gas	38590.0	automatic	sedan	silver	0	2018-12-06	

	days_listed	day_of_week_posted	month_posted	year_posted	\
30	7	0	11	2018	
37	1	6	3	2019	
57	2	3	12	2018	
70	3	1	10	2018	
111	2	4	3	2019	
...	
47792	6	2	5	2018	
47830	6	1	7	2018	
47858	6	6	3	2019	
47870	1	1	3	2019	
47879	4	3	12	2018	

	vehicle_age	mileage_per_year
30	3	17283.0
37	10	23400.0
57	12	15500.0
70	13	14310.0
111	1	11315.0
...
47792	5	23732.0
47830	11	12455.0
47858	3	35404.0
47870	7	20643.0
47879	0	inf

[2767 rows x 18 columns]

```
[137]: vehicles_filtered.query('0 < days_listed <= 7').describe()
```

```
[137]:
```


	price	model_year	condition	cylinders	odometer \
count	2767.000000	2767.000000	2767.000000	2767.000000	2767.000000
mean	12239.842429	2009.870618	2.659198	6.089989	115186.95627
std	9076.218664	5.729176	0.725933	1.631508	58478.01055
min	5.000000	1978.000000	0.000000	3.000000	0.000000
25%	5495.000000	2006.000000	2.000000	4.000000	73661.50000
50%	9500.000000	2011.000000	3.000000	6.000000	116510.00000
75%	16987.500000	2014.000000	3.000000	8.000000	153000.00000
max	49750.000000	2019.000000	5.000000	10.000000	300000.00000

	is_4wd	days_listed	day_of_week_posted	month_posted \
count	2767.000000	2767.000000	2767.000000	2767.000000
mean	0.491507	4.790025	2.946151	6.728948
std	0.500018	1.835321	1.989577	3.506255
min	0.000000	1.000000	0.000000	1.000000
25%	0.000000	3.000000	1.000000	4.000000
50%	0.000000	5.000000	3.000000	7.000000
75%	1.000000	6.000000	5.000000	10.000000
max	1.000000	7.000000	6.000000	12.000000

	year_posted	vehicle_age	mileage_per_year
count	2767.000000	2767.000000	2767.0
mean	2018.295266	8.424648	inf
std	0.456245	5.726536	NaN
min	2018.000000	0.000000	0.0
25%	2018.000000	4.000000	11129.0
50%	2018.000000	7.000000	15657.0
75%	2019.000000	12.000000	22000.0
max	2019.000000	40.000000	inf

There are 2767 car ads which are gone within a week. Let's check out the characteristics of these cars in terms of 'price', 'model_year', 'model', 'odometer', 'condition', 'transmission' and 'color', in comparison with the whole data.

```
[138]: fig, (ax1,ax2) = plt.subplots(1,2, figsize=(10,4))
vehicles_filtered.query('0 < days_listed <= 7')['price'].plot(
    kind = 'hist', bins = 50, title = 'Popular car prices', ax=ax1)
vehicles_filtered['price'].plot(kind = 'hist', bins = 50, color = 'orange',
    title = 'Car prices',ax=ax2)
plt.tight_layout()
```




```
[139]: print('Popular car price median:',
          vehicles_filtered.query('0 < days_listed <= 7')['price'].median())
print('Popular car price mean:',
      vehicles_filtered.query('0 < days_listed <= 7')['price'].mean())
print('Car price median:', vehicles_filtered['price'].median())
print('Car price mean:', vehicles_filtered['price'].mean())
```

```
Popular car price median: 9500.0
Popular car price mean: 12239.842428623058
Car price median: 9495.0
Car price mean: 12109.173805583614
```

The price distribution of the popular cars is almost identical with that of all cars. There is not much difference in the mean and median price for popular cars and for all cars either. Price doesn't seem to be a deciding factor for car popularity!


```
[140]: fig, (ax1, ax2) = plt.subplots(1, 2, figsize=(14, 4))
vehicles_filtered.query('0 < days_listed <= 7')['model_year'].value_counts().
    .plot(
        kind = 'bar', title = 'Popular car model year', ax=ax1)
ax1.set_ylabel('Frequency')
vehicles_filtered['model_year'].value_counts().plot(
    kind = 'bar', title = 'Car model year', color = 'orange', ax=ax2)
plt.tight_layout()
```



The distribution is almost identical. The top three popular car model years are 2013, 2014, and 2012, which are the same as the top three year models among all cars, just in slightly different order.

It seems that how old a car is doesn't affect its popularity. Let's have a look at the odometer next.

```
[141]: fig, (ax1,ax2) = plt.subplots(1,2, figsize=(10,4))
vehicles_filtered.query('0 < days_listed <= 7')['odometer'].plot(
    kind = 'hist', bins = 50, title = 'Popular car odometers', ax=ax1)
vehicles_filtered['odometer'].plot(kind = 'hist', bins = 50,
    title = 'Car odometers', color = 'orange',
    ax=ax2)
plt.tight_layout()
```

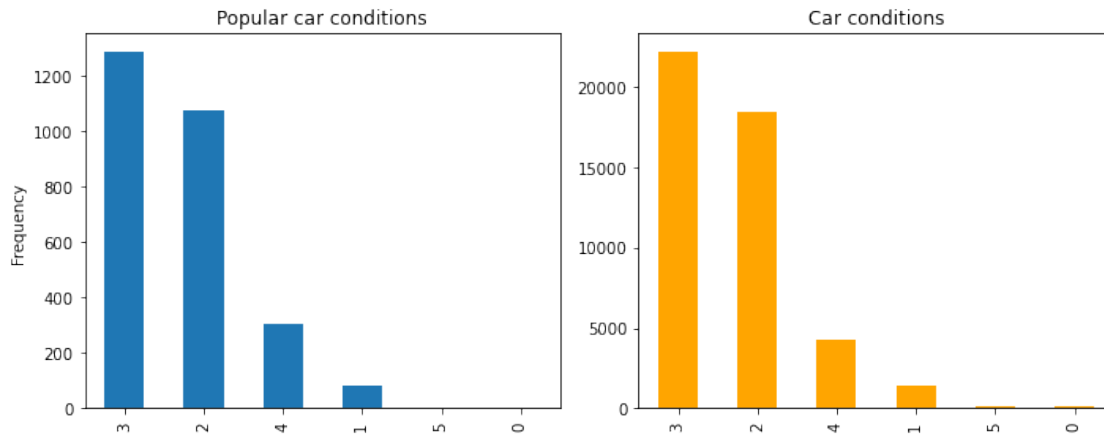


Interestingly, odometer also doesn't seem to be impacting factor. Let's have a look at condition then.

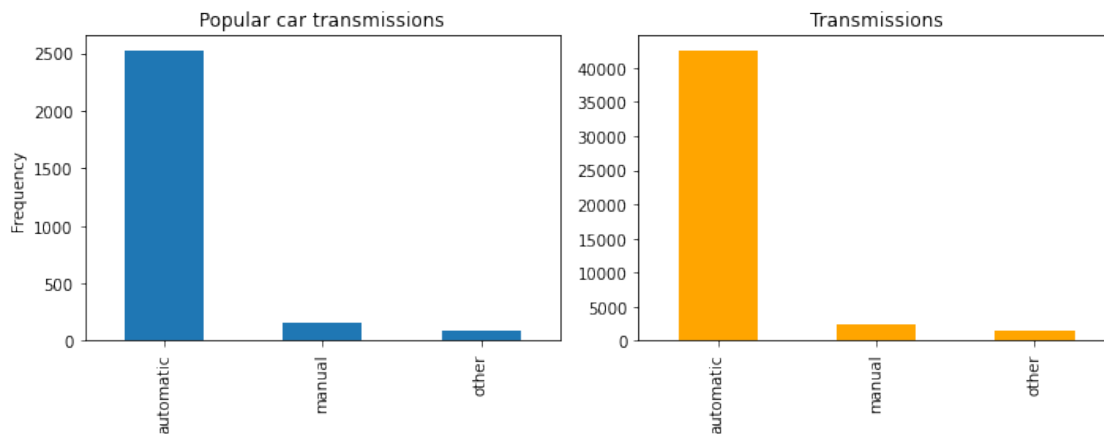
```
[142]: fig, (ax1,ax2) = plt.subplots(1,2, figsize=(10,4))
vehicles_filtered.query('0 < days_listed <= 7')['condition'].value_counts().
    plot(
    kind = 'bar', title = 'Popular car conditions', ax=ax1)
ax1.set_ylabel('Frequency')
```



```
vehicles_filtered['condition'].value_counts().plot(
    title = 'Car conditions', kind = 'bar', color = 'orange', ax=ax2)
plt.tight_layout()
```

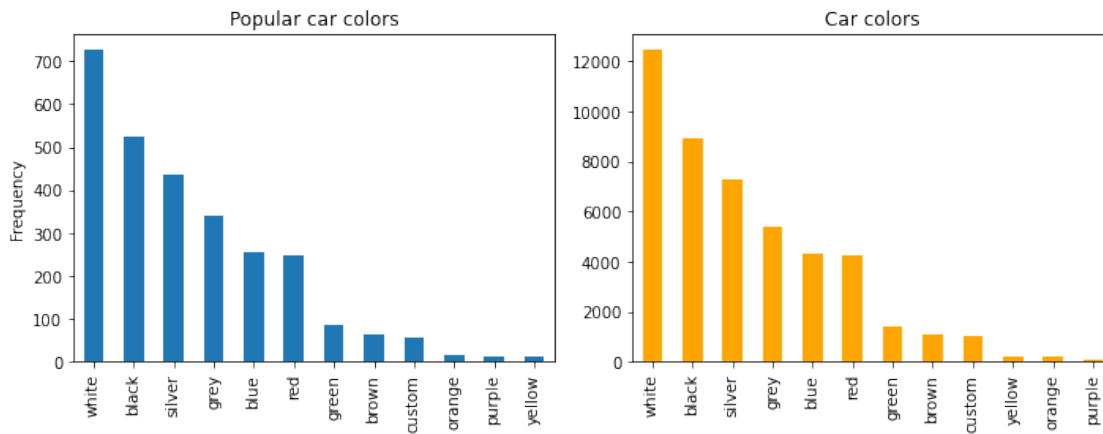


```
[143]: fig, (ax1,ax2) = plt.subplots(1,2, figsize=(10,4))
vehicles_filtered.query('0 < days_listed <= 7')['transmission'].value_counts().
    plot(
        title = 'Popular car transmissions', kind = 'bar', ax=ax1)
ax1.set_ylabel('Frequency')
vehicles_filtered['transmission'].value_counts().plot(
    kind = 'bar', title = 'Transmissions', color = 'orange', ax=ax2)
plt.tight_layout()
```



```
[145]: fig, (ax1,ax2) = plt.subplots(1,2, figsize=(10,4))
```

```
vehicles_filtered.query('0 < days_listed <= 7')['paint_color'].value_counts().
    plot(
        kind = 'bar', title= 'Popular car colors', ax=ax1)
ax1.set_ylabel('Frequency')
vehicles_filtered['paint_color'].value_counts().plot(
    kind = 'bar', color = 'orange', title = 'Car colors', ax=ax2)
plt.tight_layout()
```



All the graphing for not much! There is indeed nothing standing out for those cars that were quickly removed!

Now let's have a look at the car ad that have stayed for over 150 days.

Cars of long advertisement life

```
[146]: vehicles_filtered.query('days_listed > 150')
```

```
[146]:
```

	price	model_year		model	condition	cylinders	\
49	3800	2012		ford focus	2	4	
83	18800	2015	chevrolet	camaro lt coupe 2d	2	6	
211	8795	2014		honda civic	3	4	
647	26995	2016		chevrolet silverado	4	8	
797	8595	2014		dodge charger	4	6	
...	
47338	9995	2012		toyota tacoma	3	6	
47465	8495	2013		hyundai elantra	2	4	
47711	3500	2005		toyota camry	3	4	
47864	1200	2005		volkswagen jetta	1	5	
47877	7995	2011		chevrolet equinox	4	4	

	fuel	odometer	transmission	type	paint_color	is_4wd	date_posted	\
49	gas	130323.0	automatic	sedan	black	0	2018-11-29	
83	gas	33926.0	other	coupe	grey	0	2019-01-16	

211	gas	85452.0	automatic	sedan	grey	0	2018-09-11
647	gas	36645.0	automatic	pickup	white	1	2018-09-01
797	gas	100004.0	automatic	sedan	blue	0	2018-10-14
...
47338	gas	172695.0	automatic	truck	grey	0	2018-09-19
47465	gas	55262.0	automatic	sedan	blue	0	2018-06-30
47711	gas	208299.0	automatic	sedan	green	0	2018-06-07
47864	gas	185000.0	automatic	sedan	grey	0	2018-10-10
47877	gas	111088.0	automatic	SUV	black	0	2019-04-01

	days_listed	day_of_week_posted	month_posted	year_posted	\
49	261	3	11	2018	
83	152	2	1	2019	
211	164	1	9	2018	
647	152	5	9	2018	
797	154	6	10	2018	
...
47338	162	2	9	2018	
47465	158	5	6	2018	
47711	159	3	6	2018	
47864	158	2	10	2018	
47877	175	0	4	2019	

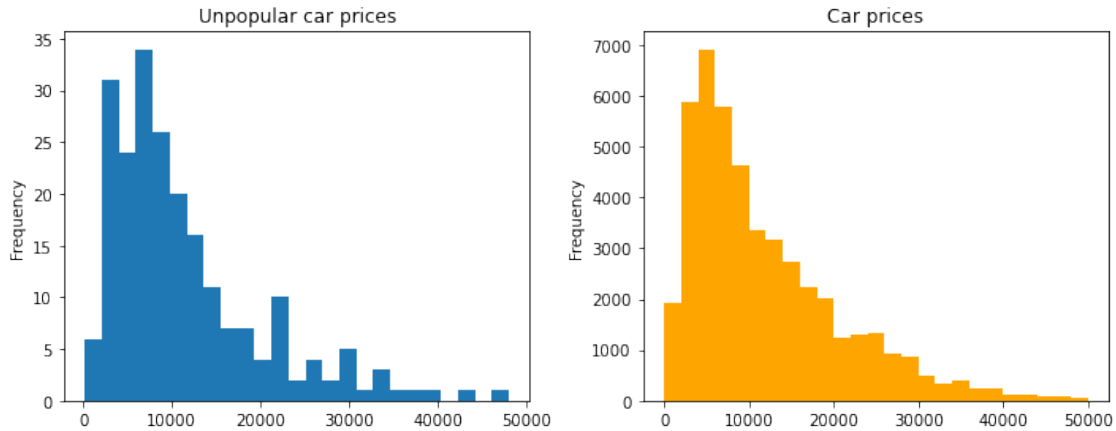
	vehicle_age	mileage_per_year
49	6	21721.0
83	4	8482.0
211	4	21363.0
647	2	18323.0
797	4	25001.0
...
47338	6	28783.0
47465	5	11053.0
47711	13	16023.0
47864	13	14231.0
47877	8	13886.0

[218 rows x 18 columns]

There are 218 car ads fitting the criterion.

Now let's compare the parameters between those cars and all cars.

```
[147]: fig, (ax1,ax2) = plt.subplots(1,2, figsize=(10,4))
vehicles_filtered.query('days_listed >150')['price'].plot(
    kind = 'hist', bins = 25, title = 'Unpopular car prices', ax=ax1)
vehicles_filtered['price'].plot(kind = 'hist', bins = 25, color = 'orange',
    title = 'Car prices', ax=ax2)
plt.tight_layout()
```




Compared to the price distribution for all cars, there seems to be a peak for not so popular car prices around 3000, then again around 6000. Let's take a look at the mean and median prices for both groups.

```
[148]: print('Unpopular car price median:',
          vehicles_filtered.query('days_listed > 150')['price'].median())
print('Unpopular car price mean:',
      vehicles_filtered.query('days_listed > 150')['price'].mean())
print('Car price median:', vehicles_filtered['price'].median())
print('Car price mean:', vehicles_filtered['price'].mean())
```

```
Unpopular car price median: 8745.0
Unpopular car price mean: 11441.188073394496
Car price median: 9495.0
Car price mean: 12109.173805583614
```

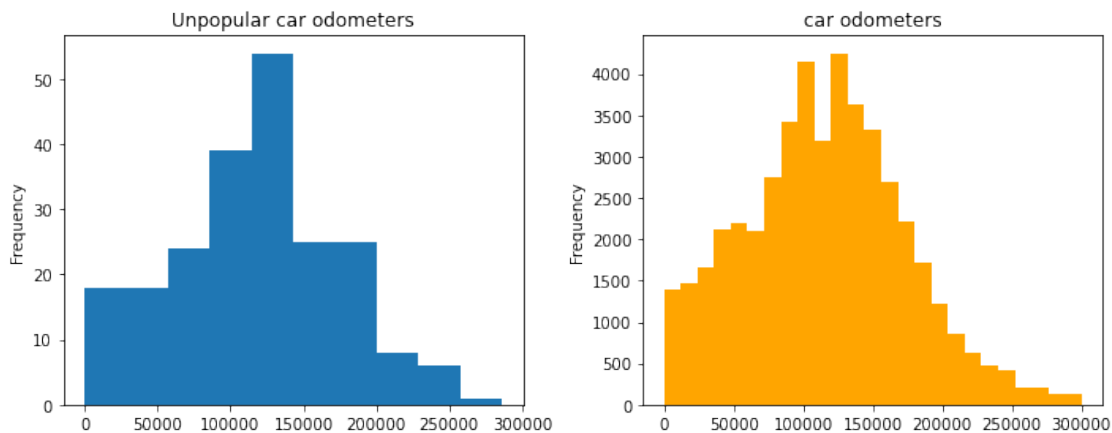
While the price distributions for both groups are similar, the unpopular cars have a slightly lower median and mean prices compared with these for all the cars.

```
[149]: fig, (ax1, ax2) = plt.subplots(1, 2, figsize=(14, 4))
vehicles_filtered.query('days_listed > 150')['model_year'].value_counts().plot(
    kind = 'bar', title = 'Unpopular car model year', ax=ax1)
ax1.set_ylabel('Frequency')
vehicles_filtered['model_year'].value_counts().plot(
    kind = 'bar', color = 'orange', title = 'Car model year', ax=ax2)
plt.tight_layout()
```



The top 5 unpopular car model years are 2011, 2013, 2008, 2012, and 2009, slightly different from the top 5 car model years: 2013, 2012, 2014, 2011, and 2015.

```
[150]: fig, (ax1,ax2) = plt.subplots(1,2, figsize=(10,4))
vehicles_filtered.query('days_listed > 150')['odometer'].plot(
    kind = 'hist', bins = 10, title = 'Unpopular car odometers', ax=ax1)
vehicles_filtered['odometer'].plot(kind = 'hist', bins = 25,
    title = 'car odometers', color = 'orange',
    ax=ax2)
plt.tight_layout()
```




```
[151]: print('Unpopular car odometer median:',
vehicles_filtered.query('days_listed > 150')['odometer'].median())
print('Unpopular car odometer mean:',
vehicles_filtered.query('days_listed > 150')['odometer'].mean())
print('Car odometer median:', vehicles_filtered['odometer'].median())
print('Car odometer mean:', vehicles_filtered['odometer'].mean())
```

Unpopular car odometer median: 118610.0
Unpopular car odometer mean: 116514.9495412844

Car odometer median: 116000.0
Car odometer mean: 115287.79697607944

As show in both graphs and the median and mean statistics, the odometers of the unpopular cars do not differ much from cars in general.

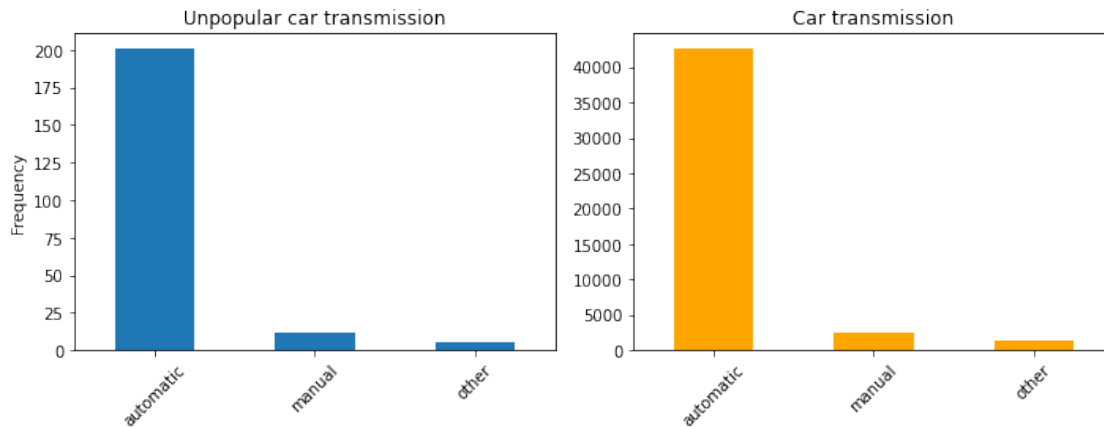
```
[152]: fig, (ax1,ax2) = plt.subplots(1,2, figsize=(10,4))
vehicles_filtered.query('days_listed >150')['condition'].value_counts().plot(
    kind = 'bar', title = 'Unpopular car condition', ax=ax1)
ax1.set_ylabel('Frequency')
vehicles_filtered['condition'].value_counts().plot(
    kind = 'bar', color = 'orange', title = 'Car condition', ax=ax2)
plt.tight_layout()
```



The top 2 conditions are 3 and 2, excellent and good, which is the same to distribution of car conditions over all. The only difference is that there is no condition '0' in the unpopular cars, which is the worse condition of all. It's interesting to see that the car conditions doesn't impact much on the car popularity either.

```
[158]: fig, (ax1,ax2) = plt.subplots(1,2, figsize=(10,4))
vehicles_filtered.query('days_listed >150')['transmission'].value_counts().plot(
    kind = 'bar', title = 'Unpopular car transmission', ax=ax1)
ax1.set_ylabel('Frequency')
ax1.set_xticklabels(ax1.get_xticklabels(), rotation= 45)

vehicles_filtered['transmission'].value_counts().plot(
    kind = 'bar', color = 'orange', title = 'Car transmission', ax=ax2)
ax2.set_xticklabels(ax2.get_xticklabels(), rotation= 45)
plt.tight_layout()
```




Both groups show almost identical distribution for car transmissions! Let's have a look at the last variable, color!

```
[159]: fig, (ax1,ax2) = plt.subplots(1,2, figsize=(10,4))
vehicles_filtered.query('days_listed >150')['paint_color'].value_counts().plot(
    kind = 'bar', title = 'Unpopular car color', ax=ax1)
ax1.set_xticklabels(ax1.get_xticklabels(), rotation= 45)
ax1.set_ylabel('Frequency')

vehicles_filtered['paint_color'].value_counts().plot(
    kind = 'bar', color = 'orange', title = 'Car color', ax=ax2)
ax2.set_xticklabels(ax2.get_xticklabels(), rotation= 45)

plt.tight_layout()
```



White and black dominate both groups, with grey and silver swapping places.

1.0.14 Conclusion

After we sorted out the cars by the number of days they were listed, we plotted the variables for both the popular ones (those that are gone within a week), and the unpopular ones (those that have stayed over 150 days), and compared the pattern in these with all cars.

Surprisingly, or not, no particular pattern really stands out for both the popular and unpopular cars, except that there seems to be a lot of cars around the price of 3000 that stay longer on the listing than others.

Well, at least we now know that it is not likely that we can predict the length of time that a car ad stays on the list based on its characteristics, but what about price? Can we predict that? We will address this query in the last section of this report.

1.0.15 4.5 Number of ads and mean price by car type

In this section, we will analyze the number of ads and the average price for each type of vehicle. A graph will be plotted to show the dependence of the number of ads on the vehicle type. We will also find out which two types have the greatest number of ads.

Let's first use `pivot_table` to calculate the number of ads and mean price for each type of car. We will sort the table in ascending order of the number of ads.

```
[160]: vehicles_type = vehicles_filtered.pivot_table(
        index = 'type', values = 'price', aggfunc = ['count', 'mean'])
vehicles_type.columns = ('number_of_ads', 'mean_price')
vehicles_type = vehicles_type.sort_values('number_of_ads')
vehicles_type
```

```
[160]:
```


	number_of_ads	mean_price
type		
bus	24	17135.666667
offroad	201	14129.855721
other	229	10661.296943
convertible	374	12674.647059
van	573	10504.680628
hatchback	962	6957.200624
mini-van	1069	8067.755847
wagon	1423	9095.635278
coupe	1982	14268.399092
pickup	6355	15948.229583
truck	11017	16720.881184
sedan	11136	7063.282058
SUV	11184	11285.182761

Now let's plot the number and price.

```
[161]: vehicles_type.plot(style = 'o-', figsize = (10, 4))
plt.xticks(rotation=45)
plt.xticks(np.arange(len(vehicles_type.index)), vehicles_type.index)
```



```
plt.show()
```



As shown in the graph above, *SUV* tops the list in its number of ads, followed by *sedan*. Now let's have a look at what factors impact their car prices most.

1.0.16 4.6 Factors in car prices

In this section, we will first generate the correlation coefficients for price and the numerical variables including age, mileage, and condition. Then we will make a scatterplot for each to show their relationship with prices.

For categorical variables (transmission type and color), a boxplot will be made to show their relationship with price.

Correlation coefficients

```
[162]: vehicles_filtered.query('type == "SUV")[[
        'price', 'vehicle_age', 'odometer', 'condition']].corr()
```

```
[162]:
```

	price	vehicle_age	odometer	condition
price	1.000000	-0.648791	-0.612848	0.294974
vehicle_age	-0.648791	1.000000	0.632784	-0.332079
odometer	-0.612848	0.632784	1.000000	-0.347473
condition	0.294974	-0.332079	-0.347473	1.000000

```
[164]: vehicles_filtered.query('type == "sedan")[[
        'price', 'vehicle_age', 'odometer', 'condition']].corr()
```

```
[164]:
```

	price	vehicle_age	odometer	condition
price	1.000000	-0.671924	-0.608063	0.311434
vehicle_age	-0.671924	1.000000	0.629729	-0.310252

```
odometer    -0.608063    0.629729    1.000000   -0.348089
condition     0.311434   -0.310252   -0.348089    1.000000
```

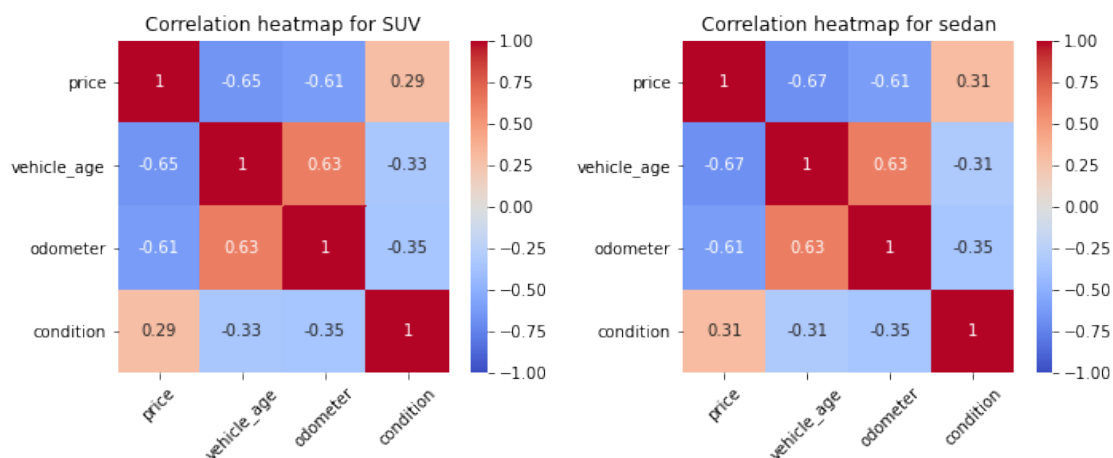
Heatmap to show the correlations

```
[173]: fig, (ax1,ax2) = plt.subplots(1,2, figsize=(10,4))

sns.heatmap(vehicles_filtered.query('type == "SUV")[[
    'price','vehicle_age', 'odometer', 'condition']].corr(),
            annot = True, fmt='.2g', vmin=-1, vmax=1,
            center= 0,cmap = 'coolwarm', square = True, ax = ax1)
ax1.set_title('Correlation heatmap for SUV')
ax1.set_xticklabels(ax1.get_xticklabels(), rotation= 45)

sns.heatmap(vehicles_filtered.query('type == "sedan")[[
    'price','vehicle_age', 'odometer', 'condition']].corr(),
            annot = True, fmt='.2g', vmin=-1, vmax=1,
            center= 0,cmap = 'coolwarm', square = True, ax = ax2)
ax2.set_title('Correlation heatmap for sedan')
ax2.set_xticklabels(ax2.get_xticklabels(), rotation= 45)

plt.tight_layout()
```

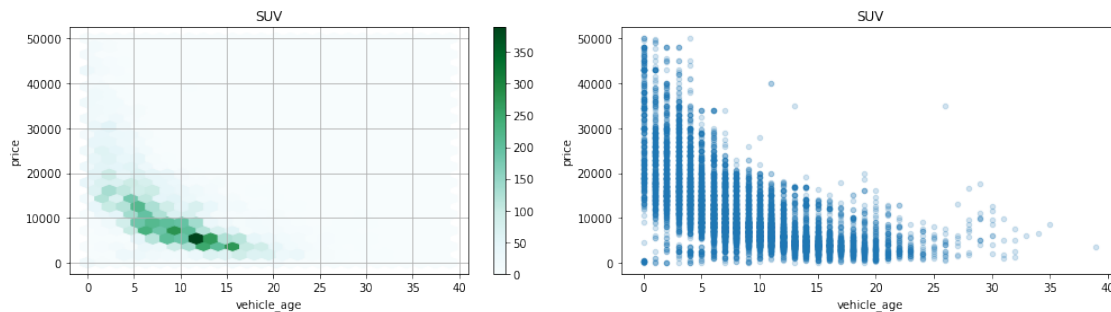


The above numbers show that, for both SUVs and sedans, car prices have relatively strong negative relationship with their age, and total mileage, whereas their relationship with car's condition is not as strong. Let's take a closer look at the scatterplot of each of these pairs.

Price and vehicle age

Give the large number of data, let's use both *hexbin* and *scatterplot* to depict the relationship among the data.


```
[174]: fig, (ax1,ax2) = plt.subplots(1,2, figsize=(14,4))
vehicles_filtered.query('type == "SUV").plot(
    x = 'vehicle_age', y = 'price', kind = 'hexbin',
    gridsize=25, sharex=False, grid=True, title = 'SUV', ax = ax1)
vehicles_filtered.query('type == "SUV").plot(
    x = 'vehicle_age', y = 'price', kind = 'scatter', title = 'SUV', alpha = 0.
    ↪2, ax = ax2)
plt.tight_layout()
```



The graphs above show that, as SUV cars get older, not only the price overall decreases, the range of prices also become smaller.

SUV of approximately 11 years and having the price around 7000 seem to be the most popular among the ads. Now let's have a look at the sedans.


```
[175]: fig, (ax1,ax2) = plt.subplots(1,2, figsize=(14,4))
vehicles_filtered.query('type == "sedan").plot(
    x = 'vehicle_age', y = 'price', kind = 'hexbin',
    gridsize=25, sharex=False, grid=True, title = 'Sedan', ax = ax1)
vehicles_filtered.query('type == "sedan").plot(
    x = 'vehicle_age', y = 'price', kind = 'scatter', title = 'Sedan', alpha = 0.
    ↪0.2, ax = ax2)
plt.tight_layout()
```




Sedans show a similar pattern as SUV, with overall cheaper prices, and younger car age.

Price and odometers

```
[176]: fig, (ax1,ax2) = plt.subplots(1,2, figsize=(14,4))
vehicles_filtered.query('type == "SUV"').plot(
    x = 'odometer', y = 'price', kind = 'hexbin',
    gridsize=25, sharex=False, grid=True, title = 'SUV', ax = ax1)
vehicles_filtered.query('type == "SUV"').plot(
    x = 'odometer', y = 'price', kind = 'scatter', title = 'SUV', alpha = 0.2,
    ax = ax2)
plt.tight_layout()
```




```
[177]: fig, (ax1,ax2) = plt.subplots(1,2, figsize=(14,4))
vehicles_filtered.query('type == "sedan"').plot(
    x = 'odometer', y = 'price', kind = 'hexbin',
    gridsize=25, sharex=False, grid=True, title = 'Sedan', ax = ax1)
vehicles_filtered.query('type == "sedan"').plot(
    x = 'odometer', y = 'price', kind = 'scatter', title = 'Sedan', alpha = 0.
    2, ax = ax2)
plt.tight_layout()
```



For both SUV and sedan cars, the graphs show a moderately strong negative correlation between prices and their total mileage. This makes sense, the older a car is, the cheaper it might sell.

Price and car conditions

```
[178]: fig, (ax1,ax2) = plt.subplots(1,2, figsize=(14,4))
vehicles_filtered.query('type == "SUV"').plot(
    x = 'condition', y = 'price', kind = 'scatter', title = 'SUV', alpha = 0.2,
    ↪ax = ax1)
vehicles_filtered.query('type == "sedan"').plot(
    x = 'condition', y = 'price', kind = 'scatter', title = 'Sedan', alpha = 0.
    ↪2, ax = ax2)
plt.tight_layout()
```



The graphs show that cars of both **SUV** and **sedan** type cluster around condition 2, 3, 4, which correspond to **good**, **excellent**, **like new**, and the prices for these car have a wide range.

Cars of condition 0 and 1, which are salvage and poor, are at the lower end of the price range. There are not as many new cars and but their prices seem to vary.

Price and transmission types

Before we plot the prices for different transmission types and colors, let's check how many ads there are in each category.

```
[179]: vehicles_filtered.query('type == "SUV"')['transmission'].value_counts()
```

```
[179]: automatic    10629
manual          468
other           87
Name: transmission, dtype: int64
```

```
[180]: vehicles_filtered.query('type == "sedan"')['transmission'].value_counts()
```

```
[180]: automatic    10518
manual         547
other          71
Name: transmission, dtype: int64
```

```
[181]: vehicles_filtered.query('type == "SUV"')['paint_color'].value_counts()
```


```
[181]: black      3020
      white      2004
      silver     1865
      grey       1254
      blue       1034
      red        836
      green      432
      brown      326
      custom     259
      orange      73
      yellow     44
      purple     37
      Name: paint_color, dtype: int64
```

```
[182]: vehicles_filtered.query('type == "sedan"')['paint_color'].value_counts()
```

```
[182]: silver     2769
      black     2259
      grey      1953
      white     1378
      blue     1189
      red       861
      custom    246
      brown     236
      green     191
      purple     25
      yellow     21
      orange      8
      Name: paint_color, dtype: int64
```

For SUVs, yellow and purple cars have less than 50 in number, and for sedans, purple, yellow, and orange has less than 50. The boxplot wouldn't work for these categories and therefore let's remove these before the plotting.

```
[183]: fig, (ax1,ax2) = plt.subplots(1,2, figsize=(14,4), sharey = True)
      vehicles_filtered.query('type == "SUV"').boxplot(
          column = 'price', by = 'transmission', ax = ax1)
      ax1.set_title('SUV')
      vehicles_filtered.query('type == "sedan"').boxplot(
          column = 'price', by = 'transmission', ax = ax2)
      ax2.set_title('Sedan')
      plt.tight_layout()
```



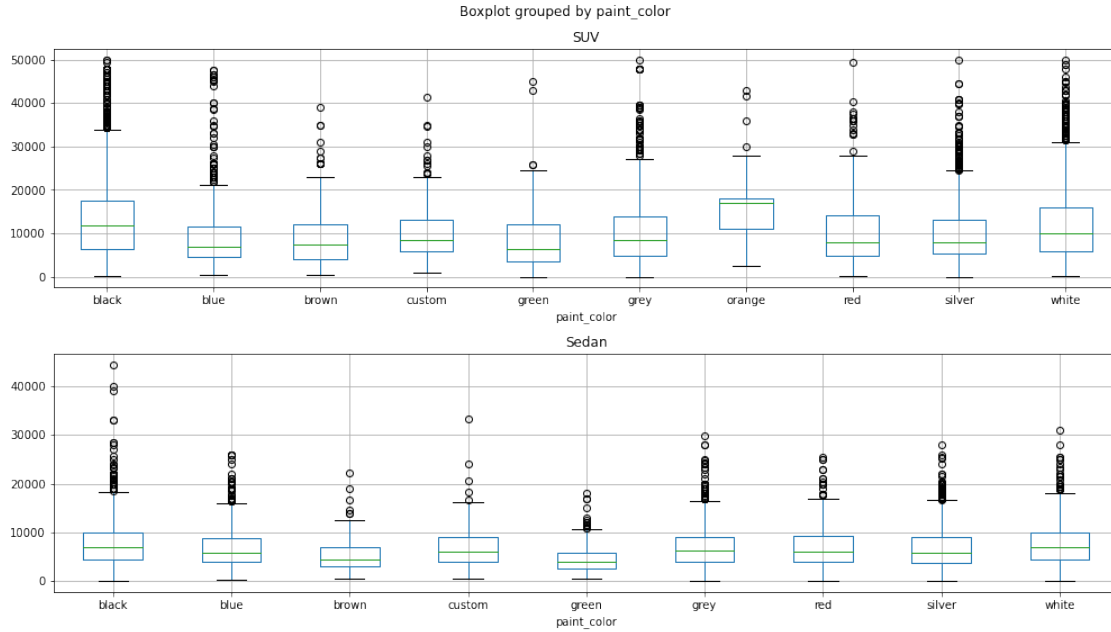
As shown in the graphs above, for the **SUV** type, manual cars have slightly higher median prices than automatic cars, whereas *other* have higher median prices than manual SUVs. It would be interesting to find out what *other* means. Automatic SUVs also have a lot of outliers towards the higher end of the price range.

In comparison, the median prices for automatic **sedans** are more expensive than manual sedans, with both transmission types having a rather dense distribution of outliers towards the higher price end. Similar to SUVs, *other* transmission type also has a higher median price than the other two types.

The IQR for sedans, shown as the boxes in the boxplots, are narrower than that for SUVs across all transmission types. This indicates that sedan prices have less dispersion than SUV prices.

```
[184]: fig, (ax1, ax2) = plt.subplots(2, 1, figsize=(14, 8))
vehicles_filtered.query(
    'type == "SUV" & paint_color not in ("purple", "yellow")').boxplot(
    column = 'price', by = 'paint_color', ax = ax1)

vehicles_filtered.query(
    'type == "sedan" & paint_color not in ("orange", "purple", "yellow")').
    boxplot(
    column = 'price', by = 'paint_color', ax = ax2)
ax1.set_title('SUV')
ax2.set_title('Sedan')
plt.tight_layout()
```



For **SUVs**, orange cars have the highest median price, followed by black and white cars. However it needs to be pointed out that there are only a total of 73 orange cars. Green SUVs have the lowest median price. Black and white SUVs also have quite densely distributed outliers over the upper whisker. Overall the car prices are right skewed.

For **sedans**, black, custom, grey, red, silver and white have relatively higher median prices than the others. Green sedans have the lowest median price. Similar to SUVs, the prices for all colors are right skewed. However, the sedans do not have as wide dispersion as SUVs, as indicated by their narrower boxes.

1.0.17 Part 5: Overall conclusion

The exploratory analysis in this report examined a dataset of free advertisements for vehicles on Crankshaft list over slightly less than a year.

The data was preprocessed by having the rows containing missing model years removed. The outliers as identified by car price, age, and mileage have also been moved to a separate dataset.

Using the preprocessed data, we examined the distribution of price, model year, cylinders, and mileage. The dominant price range seems to be around 5000, age 4 to 7 years, mileage around 125,000. The most popular cylinder categories are 6, 8, and 4, and the most popular conditions are 'good' and 'excellent'.

Examining the number of days the advertisements are listed, we discovered those that were gone with a week, and those that have stayed for over 5 months! However, a closer look at each of the parameters of those quick gone and 'permanent resident' cars doesn't really show any particular patterns.

We also found that the types of cars that have the most advertisements are SUVs and sedans, whose prices are most negatively impacted by their mileage and age.

Finally, it would also be interesting to examine the outliers, depending on what questions we are asking. However, this is beyond the scope of this report. There are also a couple of issues that might need attention of the team who provided the data, such as the missing model year values, and the 1 dollar car price.

1.0.18 Project completion checklist

- ☒ file opened
- ☒ files explored (first rows printed, info() method)
- ☒ missing values determined
- ☒ missing values filled in
- ☒ clarification of the discovered missing values provided
- ☒ data types converted
- ☒ explanation of which columns had the data types changed and why
- ☒ calculated and added to the table: day of the week, month, and year the ad was placed
- ☒ calculated and added to the table: the vehicle's age (in years) when the ad was placed
- ☒ calculated and added to the table: the vehicle's average mileage per year
- ☒ the following parameters investigated: price, vehicle's age when the ad was placed, mileage, number of cylinders, and condition
- ☒ histograms for each parameter created
- ☒ task completed: "Determine the upper limits of outliers, remove the outliers and store them in a separate DataFrame, and continue your work with the filtered data."
- ☒ task completed: "Use the filtered data to plot new histograms. Compare them with the earlier histograms (the ones that included outliers). Draw conclusions for each histogram."
- ☒ task completed: "Study how many days advertisements were displayed (days_listed). Plot a histogram. Calculate the mean and median. Describe the typical lifetime of an ad. Determine when ads were removed quickly, and when they were listed for an abnormally long time."
- ☒ task completed: "Analyze the number of ads and the average price for each type of vehicle. Plot a graph showing the dependence of the number of ads on the vehicle type. Select the two types with the greatest number of ads."
- ☒ task completed: "What factors impact the price most? Take each of the popular types you detected at the previous stage and study whether the price depends on age, mileage, condition, transmission type, and color. For categorical variables (transmission type and color), plot box-and-whisker charts, and create scatterplots for the rest. When analyzing categorical variables, note that the categories must have at least 50 ads; otherwise, their parameters won't be valid for analysis."
- ☒ each stage has a conclusion
- ☒ overall conclusion drawn